

1. (Ex. 4 Pg. 168)

a) To predict 10% of x , x must be an element of $[0.05, 0.95]$, which can be rewritten as $[x-0.05, x+0.05]$. However if $x < 0.05$, then we use observations in interval $[0, x+0.05]$ which represents $(100x+5)\%$ and if $x > 0.95$, then we use observations in interval $[x-0.05, 100]$, which represents $(105-100x)\%$.

$$\int_0^{0.05} (100x + 5) dx + \int_{0.05}^{0.95} 10 dx + \int_{0.95}^1 (105 - 100x) dx = 9.75$$

b) $0.975 * 0.975 = 0.00950625\%$ assuming the two predictors are independent

c) Assuming the predictors are independent, the fraction of available observations is $0.975^{100}\%$.

d) As p predictors increase, the fraction of available observations we use for prediction approaches 0, which means nothing to use to train our data.

e) For $p = 1, 2, 100$, we have length $= 0.1, 0.1^{1/2}, 0.1^{1/100}$

2. (Ex. 6 Pg. 170)

a) Using the logistic model, $p(X) = \frac{e^{-6+0.05X_1+X_2}}{1+e^{-6+0.05X_1+X_2}} = 0.3775$

Plug in $X_1 = 40$ hrs. and $X_2 = 3.75$ GPA and you should get the above answer

We use this model, because $p(X) = \Pr(Y = 1 | X)$, where Y is getting an A and $X = (X_1, X_2)$

b) Given a fixed GPA of 3.5, $\frac{e^{-6+0.05X_1+3.5}}{1+e^{-6+0.05X_1+3.5}} = 0.5$

Now we solve for X_1 by rewriting the equation as

$$e^{-6+0.05X_1+3.5} = 0.5 * 1 + 0.5 * e^{-6+0.05X_1+3.5}$$

$$0.5 * e^{-6+0.05X_1+3.5} = 0.5$$

$$e^{-6+0.05X_1+3.5} = 1$$

$$\log(e^{-6+0.05X_1+3.5}) = \log(1) \rightarrow -6 + 0.05 * X_1 + 3.5 = 0$$

$$0.05 * X_1 = 2.5 \rightarrow X_1 = 50$$

3. (Ex. 8 Pg. 170)

For K -nearest neighbors with $K = 1$, we have a training error rate of 0%, because $P(Y = j | X = x_i) = I(y_i = j)$. Recall when $y_i = j$, then $I = 0$, otherwise $I = 1$. Since $K = 1$, the training error rate is 0%, because of flexible classification methods. This implies our test error rate is 36% given the average of the test and training was 18%. The test error was greater than the logistic regression 30% error rate, therefore it is better to choose the logistic regression.

4. (Ex. 10 Pg. 171 a, b, c, d)

a)

Sam Lin
Stats 202
Hw. 2

> summary(Weekly)

Year	Lag1	Lag2	Lag3	Lag4
Min. :1990	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410	Median : 0.2380
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472	Mean : 0.1458
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260


Lag5	Volume	Today	Direction
Min. : -18.1950	Min. : 0.08747	Min. : -18.1950	Down:484
1st Qu.: -1.1660	1st Qu.:0.33202	1st Qu.: -1.1540	Up :605
Median : 0.2340	Median :1.00268	Median : 0.2410	
Mean : 0.1399	Mean :1.57462	Mean : 0.1499	
3rd Qu.: 1.4050	3rd Qu.:2.05373	3rd Qu.: 1.4050	
Max. : 12.0260	Max. :9.32821	Max. : 12.0260	

From this correlation function, we can tell year and volume have the biggest correlation, while correlations between Lags are near zero.

> summary(Weekly)

Year	Lag1	Lag2	Lag3	Lag4
Min. :1990	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410	Median : 0.2380
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472	Mean : 0.1458
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260

Lag5	Volume	Today	Direction
Min. : -18.1950	Min. : 0.08747	Min. : -18.1950	Down:484
1st Qu.: -1.1660	1st Qu.:0.33202	1st Qu.: -1.1540	Up :605
Median : 0.2340	Median :1.00268	Median : 0.2410	
Mean : 0.1399	Mean :1.57462	Mean : 0.1499	
3rd Qu.: 1.4050	3rd Qu.:2.05373	3rd Qu.: 1.4050	
Max. : 12.0260	Max. :9.32821	Max. : 12.0260	



b) Must pass argument family = binomial to run a logistic regression

Sam Lin

Stats 202

Hw. 2

```
> glm.fit <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = Weekly, family = binomial)
> summary(glm.fit)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = binomial, data = Weekly)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6949	-1.2565	0.9913	1.0849	1.4579

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1486.4 on 1082 degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

c) Matrix gives $(54+557)/1089 = 0.5651$. $1-0.5651 = 0.4389$ or 43.89% training error rate. When the market goes up, the model is right $557/(557+48) = 0.921$ or 92.1% of the time. When the market goes down, the model is right $54/(54+430) = 0.11157$ or 11.16% of the time.

```
> prob<-predict(glm.fit, type = "response") #predict function probabilities
> pred.glm <- rep("Down", length(prob)) #replicates elements with length
> pred.glm[prob > 0.5] <- "Up"
> table(pred.glm, Direction)
```

	Direction	
pred.glm	Down	Up
Down	54	48
Up	430	557

Sam Lin
Stats 202
Hw. 2

d)

```
> train <- (Year < 2009) #from 1990 - 2008
> glm.fit <- glm(Direction ~ Lag2, data = Weekly, family = binomial, subset = train)
> summary(glm.fit)
```

Call:

```
glm(formula = Direction ~ Lag2, family = binomial, data = Weekly,
     subset = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.536	-1.264	1.021	1.091	1.368

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.20326	0.06428	3.162	0.00157 **
Lag2	0.05810	0.02870	2.024	0.04298 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1354.7 on 984 degrees of freedom
Residual deviance: 1350.5 on 983 degrees of freedom
AIC: 1354.5

Number of Fisher Scoring iterations: 4

```
> Weekly.20092010 <- Weekly[!train, ] # for 2009 - 2010
> Direction.20092010 <- Direction[!train]
> probs2 <- predict(glm.fit, Weekly.20092010, type = "response")
> glm.fit <- rep("Down", length(probs2))
> glm.fit[probs2 > 0.5] <- "Up"
> table(glm.fit, Direction.20092010)
```

Direction.20092010

glm.fit Down Up

Down 9 5

Up 34 56

$(9+56)/104 = 62.5\%$, then $1 - 0.625 = 37.5\%$ training error rate. When the market goes up, the model is right $56/(56+5) = 91.8\%$ of the time. When the market goes down, the model is right $9/(34+9) = 20.93\%$ of the time.

5. (Ex. 5 Pg. 198)

a)

Sam Lin

Stats 202

Hw. 2

```
> attach(Default)
> set.seed(1)
> glm.fit <- glm(default ~ income + balance, family = binomial)
> summary(glm.fit)
```

Call:

```
glm(formula = default ~ income + balance, family = binomial)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
Residual deviance: 1579.0 on 9997 degrees of freedom
AIC: 1585

Number of Fisher Scoring iterations: 8

b) i. & ii. Split data and fit a model with only training data

```
> set.seed(1)
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> glm.fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> summary(glm.fit)
```

Call:

```
glm(formula = default ~ income + balance, family = "binomial",
    data = Default, subset = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.3583	-0.1268	-0.0475	-0.0165	3.8116

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.208e+01	6.658e-01	-18.148	<2e-16 ***
income	1.858e-05	7.573e-06	2.454	0.0141 *
balance	6.053e-03	3.467e-04	17.457	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1457.0 on 4999 degrees of freedom
Residual deviance: 734.4 on 4997 degrees of freedom
AIC: 740.4

Number of Fisher Scoring iterations: 8

Sam Lin

Stats 202

Hw. 2

iii. & iv.

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> #summary(fit.glm)
>
> probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0268
```

2.68% test error rate with validation set

c)

```
> probs <- predict(glm.fit, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0252
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> glm.fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> #summary(fit.glm)
>
> probs <- predict(glm.fit, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0246
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> glm.fit <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> #summary(fit.glm)
>
> probs <- predict(glm.fit, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0266
```

Validation set's test error rate varies depending on what observations are in the training set and what observations in the validation set.

d) Dummy variable student does not seem to affect the reduction of test error rate on the validation set.

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> glm.fit <- glm(default ~ income + balance + student, data = Default, family = "binomial", subset = train)
> #summary(fit.glm)
>
> probs <- predict(glm.fit, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0254
```

6. (Ex. 6 Pg. 199)

a)

Sam Lin

Stats 202

Hw. 2

```
> glm.fit <- glm(default ~ income + balance, data = Default, family = "binomial")
> summary(glm.fit)
```

Call:

```
glm(formula = default ~ income + balance, family = "binomial",
     data = Default)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom
Residual deviance: 1579.0 on 9997 degrees of freedom
AIC: 1585

Number of Fisher Scoring iterations: 8

```
> summary(glm.fit)$coefficients[,2]
(Intercept) income balance
4.347564e-01 4.985167e-06 2.273731e-04
```

The coefficients for B_0 , B_1 , B_2 are 0.4347564, 4.985167×10^{-6} , and 2.273731×10^{-4}

b)

```
> boot.fn <- function(data, index) {
+   fit <- glm(default ~ income + balance, data = data, family = "binomial", subset = index)
+   return(coef(fit))
+ }
```

c)

```
> library(boot)
> boot(Default, boot.fn, 500)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = Default, statistic = boot.fn, R = 500)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	-1.154047e+01	-5.655021e-02	4.372395e-01
t2*	2.080898e-05	-5.864198e-08	4.763111e-06
t3*	5.647103e-03	3.287761e-05	2.402848e-04

Sam Lin

Stats 202

Hw. 2


d) Notice from c) that the standard error of for B_0 , B_1 , B_2 are 0.437239, 4.763111×10^{-6} and 2.40284×10^{-4} respectively, which is close to the standard error from b) where the coefficients 0.4347564, 4.985167×10^{-6} , and 2.2733731×10^{-4} . Therefore we can conclude both estimation methods are pretty close.

7. (Ex. 8 Pg. 200)

a) In the data set, $n = 100$ observations, and $p = 2$ predictors

$$y = x - 2x^2 + \text{error}$$

b)



There is a curved relationship between x and y.

c)

i.

```
> library(boot)
> set.seed(1)
> Data <- data.frame(x,y)
> fit.glm <- glm(y~x)
> cv.glm(Data, fit.glm)$delta[1]
[1] 5.890979
```

ii.

```
> fit.glm.2 <- glm(y ~ poly(x, 2))
> cv.glm(Data, fit.glm.2)$delta[1]
[1] 1.086596
```

iii.

```
> fit.glm <- glm(y ~ poly(x, 3))
> cv.glm(Data, fit.glm)$delta[1]
[1] 1.102585
```

iv.

Sam Lin
Stats 202
Hw. 2

```
> fit.glm <- glm(y ~ poly(x, 4))
> cv.glm(Data, fit.glm)$delta[1]
[1] 1.114772
```

d) Yes, the results are identical, because leave one out cross validation evaluates n folds of a single observation

```
> set.seed(5)
> Data <- data.frame(x,y)
> fit.glm <- glm(y~x)
> cv.glm(Data, fit.glm)$delta[1]
[1] 5.890979
>
> fit.glm.2 <- glm(y ~ poly(x, 2))
> cv.glm(Data, fit.glm.2)$delta[1]
[1] 1.086596
>
> fit.glm <- glm(y ~ poly(x, 3))
> cv.glm(Data, fit.glm)$delta[1]
[1] 1.102585
>
> fit.glm <- glm(y ~ poly(x, 4))
> cv.glm(Data, fit.glm)$delta[1]
[1] 1.114772
```

e) The second fit had the smallest error, because the relationship between x and y is quadratic as we plotted in b.

f) P-values of linear and quadratic were more significant compared to cubic and quartic, which matches our cross-validation results

```
Call:
glm(formula = y ~ poly(x, 4))

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.8914  -0.5244   0.0749   0.5932   2.7796

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -1.8277      0.1041  -17.549  <2e-16 ***
poly(x, 4)1    2.3164      1.0415   2.224   0.0285 *
poly(x, 4)2  -21.0586      1.0415 -20.220  <2e-16 ***
poly(x, 4)3   -0.3048      1.0415  -0.293   0.7704
poly(x, 4)4   -0.4926      1.0415  -0.473   0.6373
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 1.084654)

Null deviance: 552.21  on 99  degrees of freedom
Residual deviance: 103.04  on 95  degrees of freedom
AIC: 298.78

Number of Fisher Scoring iterations: 2
```

8. (Ex. 9 Pg. 200)

a)

```
> muest <- mean(medv)
> muest
[1] 22.53281
```

b)

```
> seest <- sd(medv) / sqrt(dim(Boston)[1])
> seest
[1] 0.4088611
```

c)

```
> boot.fn <- function(data, index) {
+   mu <- mean(data[index])
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	22.53281	-0.01224071	0.4122534

Bootstrap estimated standard error is 0.4122534, which is very close to b

d)

```
> t.test(medv)
```

One Sample t-test

data: medv

t = 55.111, df = 505, p-value < 2.2e-16

alternative hypothesis: true mean is not equal to 0

95 percent confidence interval:

21.72953 23.33608

sample estimates:

mean of x

22.53281

T-test 95% confidence interval: [21.73, 23.33] and bootstrap interval: [21.71, 23.35] are very close.

```
> #22.53 = mean of x and 0.4119 is the standard error
```

```
> confI <- c(22.53 - 2 *0.4119, 22.53 + 2*0.4119)
```

```
> confI
```

```
[1] 21.7062 23.3538
```

e)

```
> medest <- median(medv)
> medest
[1] 21.2
```

f)

```
> boot.fn <- function(data, index) {
+   mu <- median(data[index])
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	21.2	-0.0025	0.374358

Here we see an estimated 21.2 median value, which is equivalent to the value from e. Also with standard error 0.374 which is small relative to our value.

g)

```
> percentest <- quantile(medv, c(0.1))
> percentest
10%
12.75
```

h) Similar to h, here we see an estimated tenth percentile value of 12.75, which matches our result from g with standard error approximately 0.5, but small enough to not affect our value.

```
> boot.fn <- function(data, index) {
+   mu <- quantile(data[index], c(0.1))
+   return (mu)
+ }
> boot(medv, boot.fn, 1000)
```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:

```
boot(data = medv, statistic = boot.fn, R = 1000)
```

Bootstrap Statistics :

	original	bias	std. error
t1*	12.75	0.0261	0.4912231