

Sam Lin  
Hwk. 4  
Stats 202  
Chapter 8, Exercise 10 (p. 334)

a)

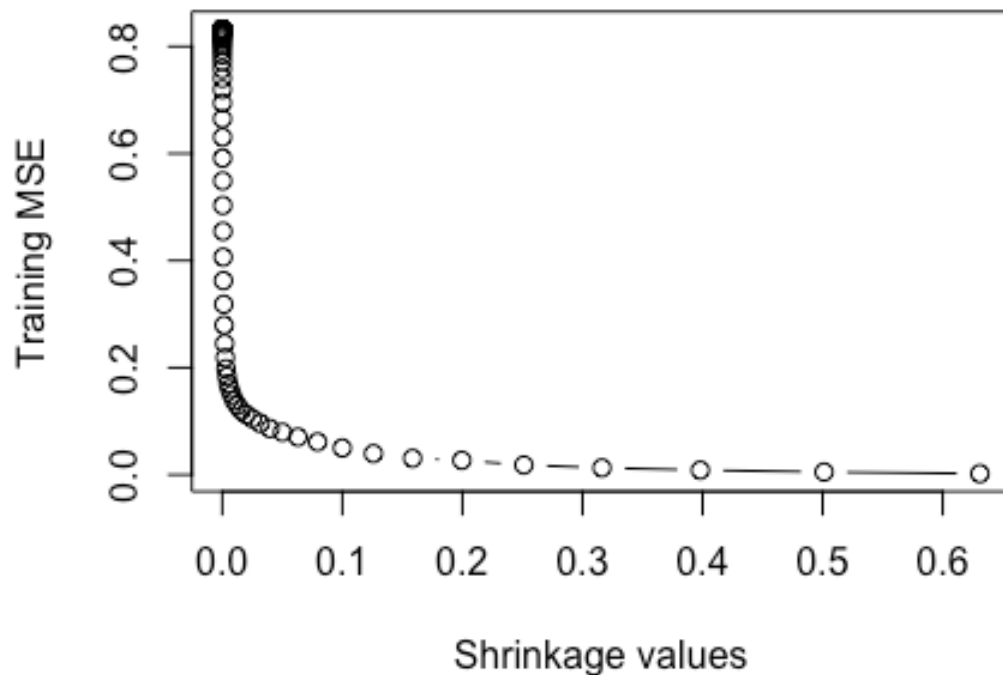
```
> Hitters <- na.omit(Hitters)
> Hitters$Salary <- log(Hitters$Salary)
```

b)

```
> train <- 1:200
> Hitters.train <- Hitters[train, ]
> Hitters.test <- Hitters[-train, ]
```

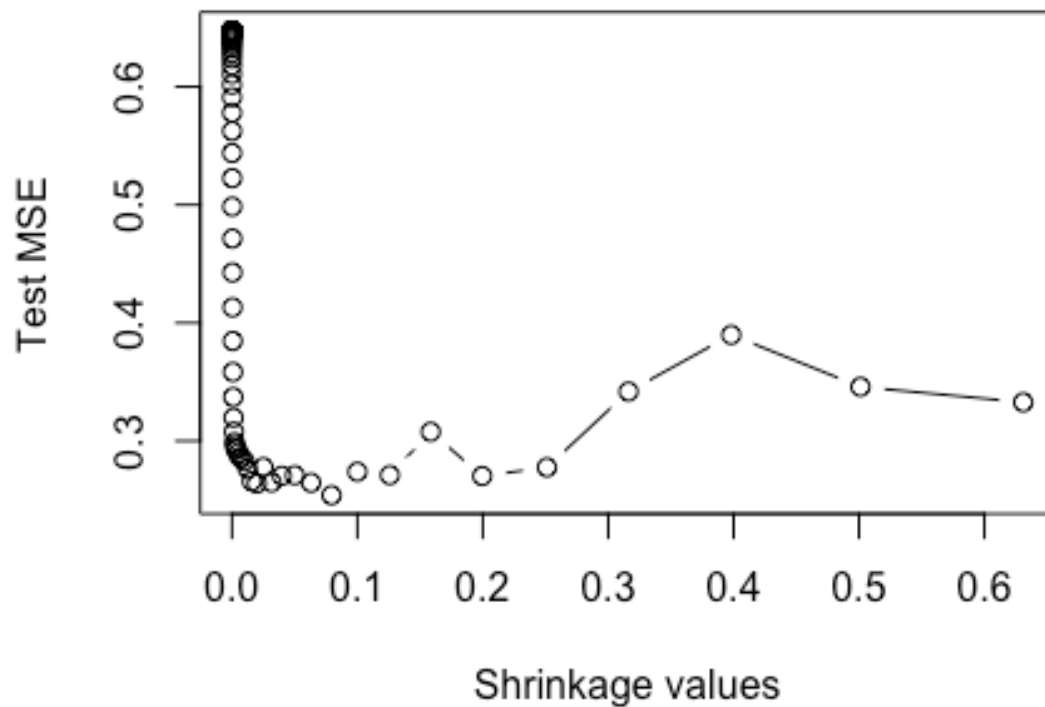
c)

```
> pows <- seq(-10, -0.2, by = 0.1)
> lambdas <- 10^pows
> train.err <- rep(NA, length(lambdas))
> for (i in 1:length(lambdas)) {
+   boost.hitters <- gbm(Salary ~ ., data = Hitters.train, distribution = "gaussian", n.trees = 1000, shrinkage = lambdas[i])
+   pred.train <- predict(boost.hitters, Hitters.train, n.trees = 1000)
+   train.err[i] <- mean((pred.train - Hitters.train$Salary)^2)
+ }
> plot(lambdas, train.err, type = "b", xlab = "Shrinkage values", ylab = "Training MSE")
```



d)

```
> for (i in 1:length(lambdas)) {
+   boost.hitters <- gbm(Salary ~ ., data = Hitters.train, distribution = "gaussian", n.trees = 1000, shrinkage = lambdas[i])
+   yhat <- predict(boost.hitters, Hitters.test, n.trees = 1000)
+   test.err[i] <- mean((yhat - Hitters.test$Salary)^2)
+ }
> plot(lambdas, test.err, type = "b", xlab = "Shrinkage values", ylab = "Test MSE")
```



```
> min(test.err)
[1] 0.2540265
> lambdas[which.min(test.err)]
[1] 0.07943282
```

Test MSE for boosting is 0.25 and  $\lambda = 0.08$

e)

```
> fit1 <- lm(Salary ~ ., data = Hitters.train)
> pred1 <- predict(fit1, Hitters.test)
> mean((pred1 - Hitters.test$Salary)^2)
[1] 0.4917959
> x <- model.matrix(Salary ~ ., data = Hitters.train)
> x.test <- model.matrix(Salary ~ ., data = Hitters.test)
> y <- Hitters.train$Salary
> fit2 <- glmnet(x, y, alpha = 0)
> pred2 <- predict(fit2, s = 0.01, newx = x.test)
> mean((pred2 - Hitters.test$Salary)^2)
[1] 0.4570283
```

Test MSE for boosting (0.25) is lower than test MSE for linear and ridge regression (0.49, 0.45 respectively)

f)

Sam Lin

Hwk. 4

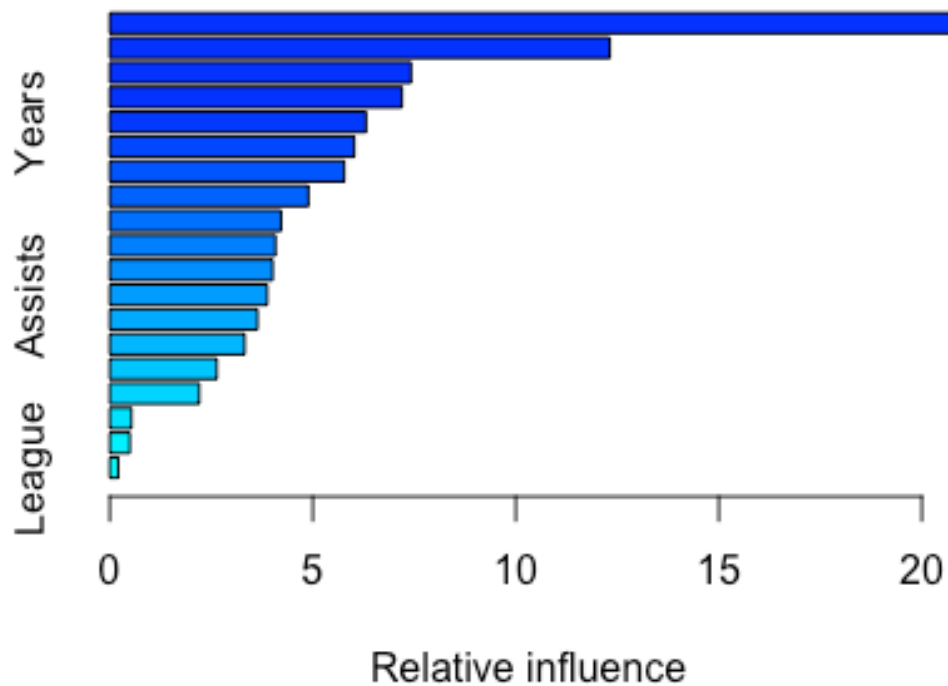
Stats 202

```
> boost.hitters <- gbm(Salary ~ ., data = Hitters.train, distribution = "gaussian", n.trees = 1000, shrinkage = lambdas[which.min(test.err)])  
> summary(boost.hitters)
```

	var	rel.inf
CatBat	CatBat	20.8404970
CRBI	CRBI	12.3158959
Walks	Walks	7.4186037
PutOuts	PutOuts	7.1958539
Years	Years	6.3104535
CWalks	CWalks	6.0221656
CHmRun	CHmRun	5.7759763
CHits	CHits	4.8914360
AtBat	AtBat	4.2187460
RBI	RBI	4.0812410
Hits	Hits	4.0117255
Assists	Assists	3.8786634
HmRun	HmRun	3.6386178
CRuns	CRuns	3.3230296
Errors	Errors	2.6369128
Runs	Runs	2.2048386
Division	Division	0.5347342
NewLeague	NewLeague	0.4943540
League	League	0.2062551

Catbat is the most important by far (rel. inf is highest)

```
> bag.hitters <- randomForest(Salary ~ ., data = Hitters.train, mtry = 19, ntree = 500)  
> yhat.bag <- predict(bag.hitters, newdata = Hitters.test)  
> mean((yhat.bag - Hitters.test$Salary)^2)  
[1] 0.2313593
```



g)

```
> bag.hitters <- randomForest(Salary ~ ., data = Hitters.train, mtry = 19, ntree = 500)  
> yhat.bag <- predict(bag.hitters, newdata = Hitters.test)  
> mean((yhat.bag - Hitters.test$Salary)^2)  
[1] 0.2313593
```

Test MSE is 0.23, which is lower than test MSE for boosting (0.25)

Sam Lin

Hwk. 4

Stats 202

Chapter 8, Exercise 11 (p. 335)

a)

```
> train <- 1:1000  
> Caravan$Purchase <- ifelse(Caravan$Purchase == "Yes", 1, 0)  
> Caravan.train <- Caravan[train, ]  
> Caravan.test <- Caravan[-train, ]
```

b)

```
> boost.caravan <- gbm(Purchase ~ ., data = Caravan.train, distribution = "gaussian", n.trees = 1000, shrinkage = 0.01)
```

c)

Sam Lin

Hwk. 4

Stats 202

> summary(boost.caravan)

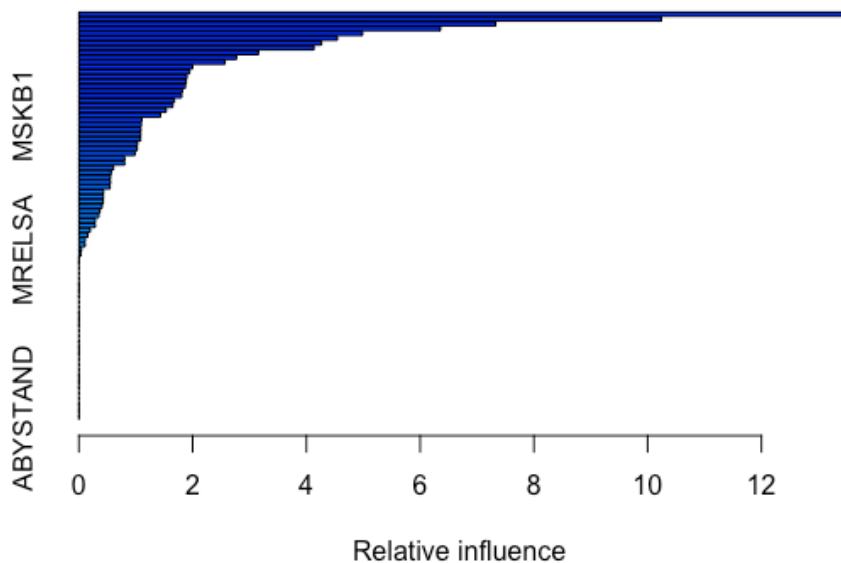
	var	rel.inf
PPERSAUT	PPERSAUT	13.51824557
MKOOKLA	MKOOKLA	10.24062778
MOPLH00G	MOPLH00G	7.32689780
MBERMIDD	MBERMIDD	6.35820558
PBRAND	PBRAND	4.98826360
ABRAND	ABRAND	4.54504653
MGODGE	MGODGE	4.26496875
MINK3045	MINK3045	4.13253907
PWAPART	PWAPART	3.15612877
MAUT1	MAUT1	2.76929763
MOSTYPE	MOSTYPE	2.56937935
MAUT2	MAUT2	1.99879666
MSKA	MSKA	1.94618539
MBERARBG	MBERARBG	1.89917331
PBYSTAND	PBYSTAND	1.88591514
MINKGEM	MINKGEM	1.87131472
MGODOV	MGODOV	1.81673309
MGODPR	MGODPR	1.80814745
MFWEKIND	MFWEKIND	1.67884570
MSKC	MSKC	1.65075962
MBERHOOG	MBERHOOG	1.53559951
MSKB1	MSKB1	1.43339514
MOPLMIDD	MOPLMIDD	1.10617074
MHHUUR	MHHUUR	1.09608784
MRELGE	MRELGE	1.09039794
MINK7512	MINK7512	1.08772012
MZFONDS	MZFONDS	1.08427551
MGODRK	MGODRK	1.03126657
MINK4575	MINK4575	1.02492795
MZPART	MZPART	0.98536712
MRELOV	MRELOV	0.80356854
MFGEKIND	MFGEKIND	0.80335689
MBERARBO	MBERARBO	0.60909852
APERSAUT	APERSAUT	0.56707821
MGEMOMV	MGEMOMV	0.55589456
MOSHOOFD	MOSHOOFD	0.55498375
MAUT0	MAUT0	0.54748481
PMOTSCO	PMOTSCO	0.43362597
MSKB2	MSKB2	0.43075446
MSKD	MSKD	0.42751490
MINK123M	MINK123M	0.40920707
MINKM30	MINKM30	0.36996576

Sam Lin

Hwk. 4

Stats 202

The top 3 most important variables are PPERSAUT, MKOOPKLA, and MOPLHOOG (due to rel. inf)



c)

```
> probs.test <- predict(boost.caravan, Caravan.test, n.trees = 1000, type = "response")
> pred.test <- ifelse(probs.test > 0.2, 1, 0)
> table(Caravan.test$Purchase, pred.test)
pred.test
  0    1
0 4493  40
1   278  11
```

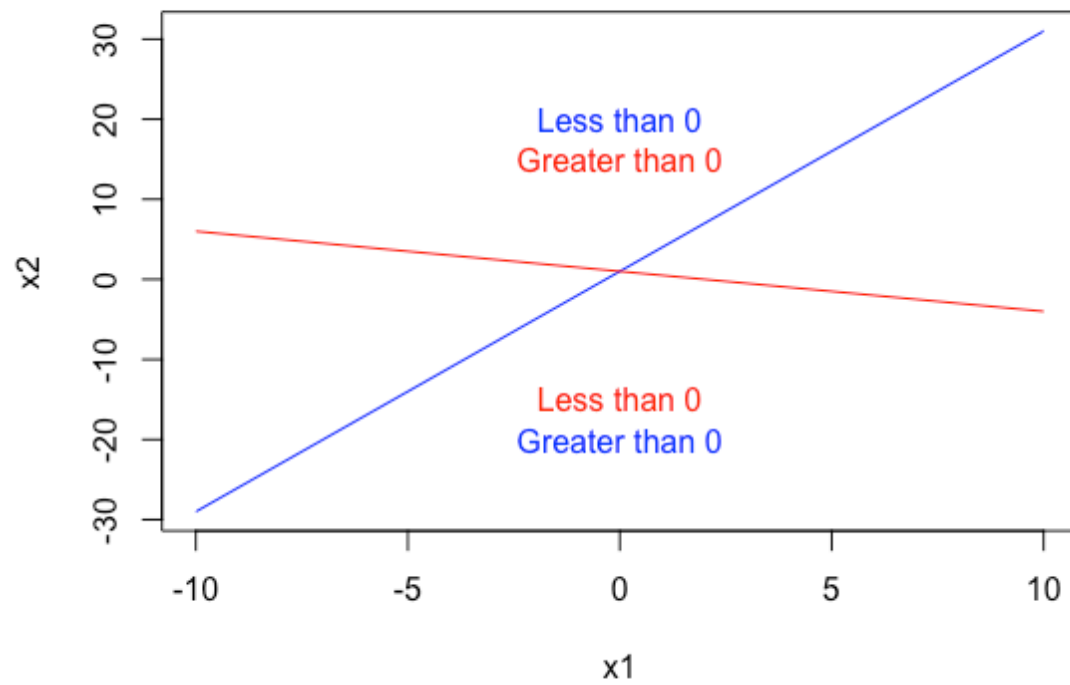
For boosting, the percentage of people buying is  $11/(11+40) = 0.2156$

```
> logit.caravan <- glm(Purchase ~ ., data = Caravan.train, family = "binomial")
Warning message:
glm.fit: fitted probabilities numerically 0 or 1 occurred
> probs.test2 <- predict(logit.caravan, Caravan.test, type = "response")
Warning message:
In predict.lm(object, newdata, se.fit, scale = 1, type = ifelse(type == :
  prediction from a rank-deficient fit may be misleading
> pred.test2 <- ifelse(probs.test > 0.2, 1, 0)
> table(Caravan.test$Purchase, pred.test2)
pred.test2
  0    1
0 4493  40
1   278  11
```

For logistic regression, the percentage of people buying is  $11/(11+40) = 0.2156$

Chapter 9, Exercise 1 (p. 368)

a)



Chapter 9, Exercise 8 (p. 371)

a)

```
> train <- sample(nrow(OJ), 800)
> OJ.train <- OJ[train, ]
> OJ.test <- OJ[-train, ]
```

Sam Lin

Hwk. 4

Stats 202

b)

```
> svm.linear <- svm(Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.01)
> summary(svm.linear)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.01)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: linear
cost: 0.01
gamma: 0.05555556
```

Number of Support Vectors: 432

```
( 215 217 )
```

Number of Classes: 2

Levels:

```
CH MM
```

432 support vectors out of 800 observations/training points. 215 support vectors belong to CH and 217 belong to MM.

c)

```
> train.pred <- predict(svm.linear, OJ.train)
> table(OJ.train$Purchase, train.pred)
```

```
train.pred
  CH  MM
CH 439  55
MM  78 228
```

Training error rate is  $(78+55)/(439+78+55+228) = 133/800 = 0.16625$

```
> test.pred <- predict(svm.linear, OJ.test)
> table(OJ.test$Purchase, test.pred)
```

```
test.pred
  CH  MM
CH 141  18
MM  31  80
```

Test error rate is  $(31+18)/(141+18+31+80) = 49/270 = 0.18148$

d) Best cost is 0.1



Sam Lin

Hwk. 4

Stats 202

```
> tune.out <- tune(svm, Purchase ~ ., data = OJ.train, kernel = "linear", ranges = list(cost = 10^seq(-2, 1, by = 0.25)))  
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost  
0.1

- best performance: 0.1625

- Detailed performance results:

	cost	error	dispersion
1	0.01000000	0.17125	0.05172376
2	0.01778279	0.16500	0.05197489
3	0.03162278	0.16625	0.04604120
4	0.05623413	0.16500	0.04594683
5	0.10000000	0.16250	0.04787136
6	0.17782794	0.16250	0.04249183
7	0.31622777	0.16875	0.04379958
8	0.56234133	0.16625	0.03998698
9	1.00000000	0.16500	0.03670453
10	1.77827941	0.16625	0.03682259
11	3.16227766	0.16500	0.03717451
12	5.62341325	0.16500	0.03525699
13	10.00000000	0.16750	0.03917553

e)

```
> test.pred <- predict(svm.linear, OJ.test)  
> table(OJ.test$Purchase, test.pred)
```

test.pred		
	CH	MM
CH	140	19
MM	32	79

With best cost, the test error rate is  $(32+19)/(32+19+79+140) = 51/270 = 0.188889$

```
> svm.linear <- svm(Purchase ~ ., kernel = "linear", data = OJ.train, cost = tune.out$best.parameter$cost)  
> train.pred <- predict(svm.linear, OJ.train)  
> table(OJ.train$Purchase, train.pred)
```

train.pred		
	CH	MM
CH	438	56
MM	71	235

With best cost, training error rate is  $(71+56)/(71+56+235+438) = 127/800 = 0.15875$

f)

Sam Lin

Hwk. 4

Stats 202

```
> svm.radial <- svm(Purchase ~ ., kernel = "radial", data = OJ.train)
> summary(svm.radial)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "radial")
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1
gamma: 0.05555556
```

Number of Support Vectors: 379

```
( 188 191 )
```

Number of Classes: 2

Levels:

```
CH MM
```

```
> train.pred <- predict(svm.radial, OJ.train)
```

```
> table(OJ.train$Purchase, train.pred)
```

```
train.pred
  CH  MM
CH 455  39
MM  77 229
```

```
> test.pred <- predict(svm.radial, OJ.test)
```

```
> table(OJ.test$Purchase, test.pred)
```

```
test.pred
  CH  MM
CH 141  18
MM  28  83
```

Created 379 support vectors out of the total observations with 188 belonging to CH and 191 belonging to MM. The classifier has training error of  $(77+39)/(455+39+77+229) = 116/800 = 0.145$  and test error =  $(18+28)/(141+18+28+83) = (46/270) = 0.1704$  which is better than linear with 15% and 18% training and testing error respectively.

## Stats 202

[illegible]

### Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost

1

- best performance: 0.16625

- Detailed performance results:

cost    error   dispersion

```
1  0.01000000 0.38250 0.04533824
```

2 0.01778279 0.38250 0.04533824

3 0.03162278 0.37500 0.04894725

4 0.05623413 0.21500 0.05886661

```
5  0.10000000 0.17875 0.04860913
```

6 0.17782794 0.17875 0.05497790

```
7  0.31622777 0.17875 0.05981743
```

```
8  0.56234133 0.17250 0.05458174
```

```
9  1.00000000 0.16625 0.05001736
```

```
10  1.77827941 0.16875 0.05008673
```

```
11 3.16227766 0.17500 0.04787136
```

12 5.62341325 0.18000 0.05244044

```
> svm.radial <- svm(Purchase ~ ., kernel = "radial", data = OJ.train, cost = tune.out$best.parameter$cost)
```

```
> summary(svm.radial)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "radial", cost = tune.out$best.parameter$cost)
```

### Parameters:

SVM-Type: C-classification

SVM-Kernel: radial

cost: 1

gamma: 0.05555556

Number of Support Vectors: 379

( 188 191 )

Number of Classes: 2

### Levels:

CH MM

```
> train.pred <- predict(svm.radial, OJ.train)
```

```
> table(OJ.train$Purchase, train.pred)
```

train.pred

CH MM

CH 455 39

MM 77 229

1

Sam Lin

Hwk. 4

Stats 202

```
> test.pred <- predict(svm.radial, OJ.test)
```

```
> table(OJ.test$Purchase, test.pred)
```

	test.pred	
	CH	MM
CH	141	18
MM	28	83

Train error:  $(77+39)/(455+39+77+229) = 116/800$

Test error :  $(28+18)/(141+18+28+83) = 46/270$

Both are same as above, so we can conclude tuning does not reduce the error rates.

g)

```
> svm.poly <- svm(Purchase ~ ., kernel = "polynomial", data = OJ.train, degree = 2)
```

```
> summary(svm.poly)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "polynomial", degree = 2)
```

Parameters:

SVM-Type: C-classification  
SVM-Kernel: polynomial  
cost: 1  
degree: 2  
gamma: 0.05555556  
coef.0: 0

Number of Support Vectors: 454

( 224 230 )

Number of Classes: 2

Levels:

CH MM

Sam Lin  
Hwk. 4  
Stats 202

```
> test.pred <- predict(svm.poly, OJ.test)
> table(OJ.test$Purchase, test.pred)
      test.pred
      CH  MM
CH 149  10
MM  41  70

> train.pred <- predict(svm.poly, OJ.train)
> table(OJ.train$Purchase, train.pred)
      train.pred
      CH  MM
CH 461  33
MM 105 201
```

Polynomial kernel with default gamma creates 454 support vectors with 224 belonging to CH and 230 belonging to MM. The classifier has a test error of  $(41+10)/(149+10+70+41) = 51/270 = 0.1725$  and training error of  $(105+33)/(461+33+105+201) = 138/800 = 0.1889$ . This is the same as the linear kernel.

```
> set.seed(2)
> tune.out <- tune(svm, Purchase ~ ., data = OJ.train, kernel = "polynomial", degree = 2, ranges = list(cost = 10^seq(-2,
+                                                                    1, by = 0.25)))
> summary(tune.out)

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:
  cost
    10

- best performance: 0.18125

- Detailed performance results:
      cost error dispersion
1  0.01000000 0.38250 0.04533824
2  0.01778279 0.36750 0.04972145
3  0.03162278 0.36500 0.05458174
4  0.05623413 0.33375 0.05070681
5  0.10000000 0.32500 0.04677072
6  0.17782794 0.25875 0.05952649
7  0.31622777 0.21250 0.06123724
8  0.56234133 0.21250 0.05743354
9  1.00000000 0.19750 0.06687468
10 1.77827941 0.19375 0.05376453
11 3.16227766 0.19625 0.05653477
12 5.62341325 0.18375 0.05434266
13 10.00000000 0.18125 0.05245699
```

Sam Lin

Hwk. 4

Stats 202

```
> svm.poly <- svm(Purchase ~ ., kernel = "polynomial", degree = 2, data = OJ.train, cost = tune.out$best.parameter$cost)
> summary(svm.poly)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "polynomial", degree = 2, cost = tune.out$best.parameter$cost)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: polynomial
cost: 10
degree: 2
gamma: 0.05555556
coef.0: 0
```

Number of Support Vectors: 342

```
( 170 172 )
```

Number of Classes: 2

Levels:

```
CH MM
```

```
> train.pred <- predict(svm.poly, OJ.train)
```

```
> table(OJ.train$Purchase, train.pred)
```

```
      train.pred
      CH  MM
CH 450  44
MM  72 234
```

```
> test.pred <- predict(svm.poly, OJ.test)
```

```
> table(OJ.test$Purchase, test.pred)
```

```
      test.pred
      CH  MM
CH 140  19
MM  31  80
```

Training error rate:  $(72+44)/(450+44+72+234) = 96/800 = 0.145$

Test error rate:  $(31+19)/(140+80+31+19) = 50/270 = 0.1851$

Conclusion: tuning reduces training and training error compared to 0.1725 and 0.18889 training and test error.

h)

The radial basis kernel produces the minimum classification error on test and training errors compared to others.